

# GROUPBY MECHANICS

In [1]:

```
import numpy as np
import pandas as pd
```

In [6]:

```
#Groupby Mechanics

# The term split-apply-combine is used for describing group operations.
# In the first stage of the process, data contained in a pandas object, whether a Series, Data-Frame,
# or otherwise, is split into groups based on one or more keys that you provide.
# The splitting is performed on a particular axis of an object.
# For example, a DataFrame can be grouped on its rows (axis=0) or its columns (axis=
1).
# Once this is done, a function is applied to each group, producing a new value.
# Finally, the results of all those function applications are combined into a result object.
# The form of the resulting object will usually depend on what's being done to the data.

df = pd.DataFrame({'key1' : ['a', 'a', 'b', 'b', 'a'],
                    'key2' : ['one', 'two', 'one', 'two', 'one'],
                    'data1' : np.random.randn(5),
                    'data2' : np.random.randn(5)})

df
```

Out[6]:

|   | key1 | key2 | data1     | data2     |
|---|------|------|-----------|-----------|
| 0 | a    | one  | -0.699511 | -0.388558 |
| 1 | a    | two  | -0.345863 | 0.535094  |
| 2 | b    | one  | -0.749546 | -1.071714 |
| 3 | b    | two  | -2.196750 | 0.524351  |
| 4 | a    | one  | 0.264189  | 2.621417  |

In [7]:

```
# Suppose you wanted to compute the mean of the data1 column using the labels from key1.
# There are a number of ways to do this.
# One is to access data1 and call groupby with the column (a Series) at key1:

grouped = df['data1'].groupby(df['key1'])
grouped
```

Out[7]:

```
<pandas.core.groupby.generic.SeriesGroupBy object at 0x00000039EB6D3A48>
```

In [16]:

```
# This grouped variable is now a GroupBy object. It has not actually computed anything
# yet except for some intermediate data about the group key df['key1'].
# The idea is that this object has all of the information needed to then apply some operation to
# each of the groups.
# For example, to compute group size we can call the GroupBy's size method:

grouped.size()
```

Out[16]:

```
key1
a      3
b      2
Name: data1, dtype: int64
```

In [17]:

```
# To compute group summary statistics we can call the GroupBy's describe method:

grouped.describe()
```

Out[17]:

|      | count | mean      | std      | min       | 25%       | 50%       | 75%       | max       |
|------|-------|-----------|----------|-----------|-----------|-----------|-----------|-----------|
| key1 |       |           |          |           |           |           |           |           |
| a    | 3.0   | -0.260395 | 0.487502 | -0.699511 | -0.522687 | -0.345863 | -0.040837 | 0.264189  |
| b    | 2.0   | -1.473148 | 1.023328 | -2.196750 | -1.834949 | -1.473148 | -1.111347 | -0.749546 |

In [19]:

```
# To compute group means we can call the GroupBy's mean method:
```

```
grouped.mean()
```

Out[19]:

```
key1
a    -0.260395
b    -1.473148
Name: data1, dtype: float64
```

In [20]:

```
# Thus above, the data (a Series) has been aggregated according to the group key,
# producing a new Series that is now indexed by the unique values in the key1 column.
# The result index has the name 'key1' because the DataFrame column df['key1'] did.
# If instead we had passed multiple arrays as a list, we'd get something different:
```

```
means = df['data1'].groupby([df['key1'], df['key2']]).mean()
means
```

Out[20]:

```
key1  key2
a      one   -0.217661
      two   -0.345863
b      one   -0.749546
      two   -2.196750
Name: data1, dtype: float64
```

In [21]:

```
df.stack()
```

Out[21]:

```
0  key1      a
   key2    one
   data1 -0.699511
   data2 -0.388558
1  key1      a
   key2    two
   data1 -0.345863
   data2  0.535094
2  key1      b
   key2    one
   data1 -0.749546
   data2 -1.07171
3  key1      b
   key2    two
   data1 -2.19675
   data2  0.524351
4  key1      a
   key2    one
   data1  0.264189
   data2  2.62142
dtype: object
```

In [22]:

```
df.T
```

Out[22]:

|       | 0         | 1         | 2         | 3        | 4        |
|-------|-----------|-----------|-----------|----------|----------|
| key1  | a         | a         | b         | b        | a        |
| key2  | one       | two       | one       | two      | one      |
| data1 | -0.699511 | -0.345863 | -0.749546 | -2.19675 | 0.264189 |
| data2 | -0.388558 | 0.535094  | -1.07171  | 0.524351 | 2.62142  |

In [24]:

```
# Here we grouped the data using two keys, and the resulting Series now has a hierarchical
# index consisting of the unique pairs of keys observed:

means.unstack()
```

Out[24]:

|      | key2 | one       | two       |
|------|------|-----------|-----------|
| key1 |      |           |           |
| a    |      | -0.217661 | -0.345863 |
| b    |      | -0.749546 | -2.196750 |

In [25]:

```
# In this example, the group keys are all Series, though they could be any arrays of the r
ight length:

states = np.array(['Ohio', 'California', 'California', 'Ohio', 'Ohio'])
years = np.array([2005, 2005, 2006, 2005, 2006])
df['data1'].groupby([states, years]).mean()
```

Out[25]:

|            |      |           |
|------------|------|-----------|
| California | 2005 | -0.345863 |
|            | 2006 | -0.749546 |
| Ohio       | 2005 | -1.448130 |
|            | 2006 | 0.264189  |

Name: data1, dtype: float64

In [26]:

```
# Frequently the grouping information is found in the same DataFrame as the data you want
to work on.
# In that case, you can pass column names (whether those are strings, numbers, or othe
r Python objects)
# as the group keys:

df.groupby('key1').mean()
```

Out[26]:

|      | data1     | data2     |
|------|-----------|-----------|
| key1 |           |           |
| a    | -0.260395 | 0.922651  |
| b    | -1.473148 | -0.273681 |

In [27]:

```
df.groupby(['key1', 'key2']).mean()
```

Out[27]:

|      |      | data1     | data2     |
|------|------|-----------|-----------|
| key1 | key2 |           |           |
| a    | one  | -0.217661 | 1.116430  |
|      | two  | -0.345863 | 0.535094  |
| b    | one  | -0.749546 | -1.071714 |
|      | two  | -2.196750 | 0.524351  |

In [28]:

```
df.groupby(['key1', 'key2']).mean()
```

Out[28]:

|      |      | data1     | data2     |
|------|------|-----------|-----------|
| key1 | key2 |           |           |
| a    | one  | -0.217661 | 1.116430  |
|      | two  | -0.345863 | 0.535094  |
| b    | one  | -0.749546 | -1.071714 |
|      | two  | -2.196750 | 0.524351  |

In [29]:

```
# You may have noticed in the first case df.groupby('key1').mean() that there is no key2 column in the result.
# Because df['key2'] is not numeric data, it is said to be a nuisance column,
# which is therefore excluded from the result.
# By default, all of the numeric columns are aggregated,
# though it is possible to filter down to a subset, as you'll see soon.

# Regardless of the objective in using groupby, a generally useful GroupBy method is size,
# which returns a Series containing group sizes:

df.groupby(['key1', 'key2']).size()
```

Out[29]:

```
key1  key2
a      one    2
      two    1
b      one    1
      two    1
dtype: int64
```

In [ ]:

```
# Take note that any missing values in a group key will be excluded from the result.
```

In [30]:

```
# Iterating over the groups

# The GroupBy object supports iteration, generating a sequence of 2-tuples containing
# the group name along with the chunk of data.

for name, group in df.groupby('key1'):
    print(name)
    print(group)
```

```
a
  key1 key2    data1    data2
0    a  one -0.699511 -0.388558
1    a  two -0.345863  0.535094
4    a  one  0.264189  2.621417
b
  key1 key2    data1    data2
2    b  one -0.749546 -1.071714
3    b  two -2.196750  0.524351
```

In [31]:

```
# for multiple keys

# In the case of multiple keys, the first element in the tuple will be a tuple of key values:

for (k1, k2), group in df.groupby(['key1', 'key2']):
    print((k1, k2))
    print(group)
```

```
('a', 'one')
  key1 key2    data1    data2
0    a  one -0.699511 -0.388558
4    a  one  0.264189  2.621417
('a', 'two')
  key1 key2    data1    data2
1    a  two -0.345863  0.535094
('b', 'one')
  key1 key2    data1    data2
2    b  one -0.749546 -1.071714
('b', 'two')
  key1 key2    data1    data2
3    b  two -2.19675  0.524351
```

In [32]:

```
# computing a dict of the data pieces as a one-liner

# Of course, you can choose to do whatever you want with the pieces of data.
# A recipe you may find useful is computing a dict of the data pieces as a one-liner:

pieces = dict(list(df.groupby('key1')))
pieces
```

Out[32]:

```
{ 'a':   key1 key2    data1    data2
0    a  one -0.699511 -0.388558
1    a  two -0.345863  0.535094
4    a  one  0.264189  2.621417,
  'b':   key1 key2    data1    data2
2    b  one -0.749546 -1.071714
3    b  two -2.196750  0.524351}
```



In [33]:

```
pieces['a']
```

Out[33]:

|   | key1 | key2 | data1     | data2     |
|---|------|------|-----------|-----------|
| 0 | a    | one  | -0.699511 | -0.388558 |
| 1 | a    | two  | -0.345863 | 0.535094  |
| 4 | a    | one  | 0.264189  | 2.621417  |

In [34]:

```
pieces['b']
```

Out[34]:

|   | key1 | key2 | data1     | data2     |
|---|------|------|-----------|-----------|
| 2 | b    | one  | -0.749546 | -1.071714 |
| 3 | b    | two  | -2.196750 | 0.524351  |

In [35]:

```
# By default groupby groups on axis=0, but you can group on any of the other axes.  
# For example, we could group the columns of our example df here by dtype like so:  
df.dtypes
```

Out[35]:

```
key1      object  
key2      object  
data1    float64  
data2    float64  
dtype: object
```

In [38]:

```
grouped1 = df.groupby(df.dtypes, axis=1)  
grouped1
```

Out[38]:

```
<pandas.core.groupby.generic.DataFrameGroupBy object at 0x00000039F1B05448>
```

In [39]:

```
grouped1.size()
```

Out[39]:

```
float64    2
object     2
dtype: int64
```

In [40]:

```
# We can print out the groups like so:
```

```
for dtype, group in grouped:
    print(dtype)
    print(group)
```

```
a
0  -0.699511
1  -0.345863
4   0.264189
Name: data1, dtype: float64
b
2  -0.749546
3  -2.196750
Name: data1, dtype: float64
```

In [41]:

```
for dtype, group in grouped1:
    print(dtype)
    print(group)
```

```
float64
      data1      data2
0 -0.699511 -0.388558
1 -0.345863  0.535094
2 -0.749546 -1.071714
3 -2.196750  0.524351
4  0.264189  2.621417
object
  key1 key2
0    a  one
1    a  two
2    b  one
3    b  two
4    a  one
```

In [42]:

```
# Selecting a Column or Subset of Columns

# Indexing a GroupBy object created from a DataFrame with a column name or array
# of column names has the effect of column subsetting for aggregation.

# This means that:

df.groupby('key1')['data1']
df.groupby('key1')[['data2']]

# are syntactic sugar for:

df['data1'].groupby(df['key1'])
df[['data2']].groupby(df['key1'])
```

Out[42]:

```
<pandas.core.groupby.generic.DataFrameGroupBy object at 0x00000039F1862808>
```

In [43]:

```
# Dataframes

# Especially for large datasets, it may be desirable to aggregate only a few columns.
# For example, in the preceding dataset, to compute means for just the data2 column
# and get the result as a DataFrame, we could write:

df.groupby(['key1', 'key2'])['data2'].mean()
```

Out[43]:

|      |      | data2     |
|------|------|-----------|
| key1 | key2 |           |
| a    | one  | 1.116430  |
|      | two  | 0.535094  |
| b    | one  | -1.071714 |
|      | two  | 0.524351  |

In [45]:

```
# Series

# The object returned by this indexing operation is a grouped DataFrame if a List or
# array is passed or a grouped Series if only a single column name is passed as a scalar:

s_grouped = df.groupby(['key1', 'key2'])['data2']
print(s_grouped)
s_grouped.mean()
```

<pandas.core.groupby.generic.SeriesGroupBy object at 0x00000039F1B0DB48>

Out[45]:

```
key1  key2
a      one    1.116430
      two    0.535094
b      one   -1.071714
      two    0.524351
Name: data2, dtype: float64
```

In [48]:

```
# Grouping with Dictionaries and Series

# Grouping information may exist in a form other than an array.
# Let's consider another example DataFrame:

people = pd.DataFrame(np.random.randn(5, 5),
                      columns=['a', 'b', 'c', 'd', 'e'],
                      index=['Joe', 'Steve', 'Wes', 'Jim', 'Travis'])

people
```

Out[48]:

|        | a         | b         | c         | d         | e         |
|--------|-----------|-----------|-----------|-----------|-----------|
| Joe    | 0.061253  | -0.219182 | 0.258375  | -1.033663 | -0.642642 |
| Steve  | 0.818724  | -0.025076 | 1.151003  | 0.089388  | 0.024214  |
| Wes    | 1.197966  | -0.412310 | -1.040145 | -1.106289 | -0.742792 |
| Jim    | -0.044602 | 0.463223  | 1.254485  | 0.765563  | 0.906828  |
| Travis | -0.681018 | -0.420505 | 0.081755  | 0.707100  | 0.900215  |

In [49]:

```
people.iloc[2:3, [1, 2]] = np.nan # Add a few NA values
people
```

Out[49]:

|        | a         | b         | c        | d         | e         |
|--------|-----------|-----------|----------|-----------|-----------|
| Joe    | 0.061253  | -0.219182 | 0.258375 | -1.033663 | -0.642642 |
| Steve  | 0.818724  | -0.025076 | 1.151003 | 0.089388  | 0.024214  |
| Wes    | 1.197966  | NaN       | NaN      | -1.106289 | -0.742792 |
| Jim    | -0.044602 | 0.463223  | 1.254485 | 0.765563  | 0.906828  |
| Travis | -0.681018 | -0.420505 | 0.081755 | 0.707100  | 0.900215  |

In [52]:

```
# Now, suppose I have a group correspondence for the columns and want to sum
# together the columns by group:

mapping = {'a': 'red', 'b': 'red', 'c': 'blue', 'd': 'blue', 'e': 'red', 'f' : 'orange'}

# Now, you could construct an array from this dict to pass to groupby,
# but instead we can just pass the dict (I included the key 'f' to highlight
# that unused grouping keys are OK):

by_column = people.groupby(mapping, axis=1)

print(people)
print()
print(by_column)
```

|        | a         | b         | c        | d         | e         |
|--------|-----------|-----------|----------|-----------|-----------|
| Joe    | 0.061253  | -0.219182 | 0.258375 | -1.033663 | -0.642642 |
| Steve  | 0.818724  | -0.025076 | 1.151003 | 0.089388  | 0.024214  |
| Wes    | 1.197966  | NaN       | NaN      | -1.106289 | -0.742792 |
| Jim    | -0.044602 | 0.463223  | 1.254485 | 0.765563  | 0.906828  |
| Travis | -0.681018 | -0.420505 | 0.081755 | 0.707100  | 0.900215  |

<pandas.core.groupby.generic.DataFrameGroupBy object at 0x00000039F1B6C348>

In [56]:

```
print(by_column.sum())  
  
print()  
  
print(by_column.mean())
```

|        | blue      | red       |
|--------|-----------|-----------|
| Joe    | -0.775288 | -0.800571 |
| Steve  | 1.240391  | 0.817862  |
| Wes    | -1.106289 | 0.455174  |
| Jim    | 2.020048  | 1.325449  |
| Travis | 0.788855  | -0.201309 |

|        | blue      | red       |
|--------|-----------|-----------|
| Joe    | -0.387644 | -0.266857 |
| Steve  | 0.620196  | 0.272621  |
| Wes    | -1.106289 | 0.227587  |
| Jim    | 1.010024  | 0.441816  |
| Travis | 0.394428  | -0.067103 |

In [58]:

```
# The same functionality holds for Series, which can be viewed as a fixed-size mapping:  
  
map_series = pd.Series(mapping)  
map_series
```

Out[58]:

```
a      red  
b      red  
c      blue  
d      blue  
e      red  
f  orange  
dtype: object
```

In [59]:

```
people.groupby(map_series, axis=1).count()
```

Out[59]:

|               | blue | red |
|---------------|------|-----|
| <b>Joe</b>    | 2    | 3   |
| <b>Steve</b>  | 2    | 3   |
| <b>Wes</b>    | 1    | 2   |
| <b>Jim</b>    | 2    | 3   |
| <b>Travis</b> | 2    | 3   |

In [60]:

```
# Grouping with Functions
```

```
# Using Python functions is a more generic way of defining a group mapping compared with a dict or Series.
```

```
# Any function passed as a group key will be called once per index value,  
# with the return values being used as the group names.  
# More concretely, consider the example DataFrame from the previous section,  
# which has people's first names as index values.  
# Suppose you wanted to group by the length of the names;  
# while you could compute an array of string lengths,  
# it's simpler to just pass the len function:
```

```
people.groupby(len).sum() # grouping by the length of the names
```

Out[60]:

|   | a         | b         | c        | d         | e         |
|---|-----------|-----------|----------|-----------|-----------|
| 3 | 1.214617  | 0.244041  | 1.512859 | -1.374389 | -0.478606 |
| 5 | 0.818724  | -0.025076 | 1.151003 | 0.089388  | 0.024214  |
| 6 | -0.681018 | -0.420505 | 0.081755 | 0.707100  | 0.900215  |

In [61]:

```
# Mixing functions with arrays, dicts, or Series is not a problem as  
# everything gets converted to arrays internally:
```

```
key_list = ['one', 'one', 'one', 'two', 'two']
```

```
people.groupby([len, key_list]).min() # Mixing functions with arrays, dicts, or Series
```

Out[61]:

|   |     | a         | b         | c        | d         | e         |
|---|-----|-----------|-----------|----------|-----------|-----------|
| 3 | one | 0.061253  | -0.219182 | 0.258375 | -1.106289 | -0.742792 |
|   | two | -0.044602 | 0.463223  | 1.254485 | 0.765563  | 0.906828  |
| 5 | one | 0.818724  | -0.025076 | 1.151003 | 0.089388  | 0.024214  |
| 6 | two | -0.681018 | -0.420505 | 0.081755 | 0.707100  | 0.900215  |

In [67]:

```
# Grouping by Index Levels
```

```
# A final convenience for hierarchically indexed datasets is the ability to aggregate  
# using one of the levels of an axis index.
```

```
# Let's Look at an example:
```

```
columns = pd.MultiIndex.from_arrays([[ 'US', 'US', 'US', 'JP', 'JP'],  
                                     [1, 3, 5, 1, 3]],  
                                   names=[ 'cty', 'tenor'])  
print(columns)
```

```
print()  
print()
```

```
hier_df = pd.DataFrame(np.random.randn(4, 5), columns=columns)  
print(hier_df)
```

```
MultiIndex([('US', 1),  
            ('US', 3),  
            ('US', 5),  
            ('JP', 1),  
            ('JP', 3)],  
          names=[ 'cty', 'tenor'])
```

|       | US        |           |           | JP        |           |
|-------|-----------|-----------|-----------|-----------|-----------|
| cty   |           |           |           |           |           |
| tenor | 1         | 3         | 5         | 1         | 3         |
| 0     | 1.173735  | -0.159624 | 0.656455  | -0.989766 | -0.529011 |
| 1     | 0.923138  | 0.036186  | 1.547744  | -0.285430 | -0.267631 |
| 2     | 0.232222  | -1.022996 | 0.478855  | -1.479527 | 0.163846  |
| 3     | -0.574041 | 0.017466  | -1.255832 | -0.247150 | 0.242282  |

In [68]:

```
# To group by level, pass the level number or name using the level keyword:
```

```
hier_df.groupby(level='cty', axis=1).count()
```

Out[68]:

| cty | JP | US |
|-----|----|----|
| 0   | 2  | 3  |
| 1   | 2  | 3  |
| 2   | 2  | 3  |
| 3   | 2  | 3  |



In [71]:

```
hier_df.groupby(level='tenor', axis=1).count()
```

Out[71]:

| tenor | 1 | 3 | 5 |
|-------|---|---|---|
| 0     | 2 | 2 | 1 |
| 1     | 2 | 2 | 1 |
| 2     | 2 | 2 | 1 |
| 3     | 2 | 2 | 1 |

In [82]:

```
hier_df.groupby(level='cty', axis=1).mean()
```

Out[82]:

| cty | JP        | US        |
|-----|-----------|-----------|
| 0   | -0.759389 | 0.556855  |
| 1   | -0.276530 | 0.835690  |
| 2   | -0.657841 | -0.103973 |
| 3   | -0.002434 | -0.604136 |

In [83]:

```
hier_df.groupby(level='tenor', axis=1).mean()
```

Out[83]:

| tenor | 1         | 3         | 5         |
|-------|-----------|-----------|-----------|
| 0     | 0.091984  | -0.344317 | 0.656455  |
| 1     | 0.318854  | -0.115722 | 1.547744  |
| 2     | -0.623652 | -0.429575 | 0.478855  |
| 3     | -0.410595 | 0.129874  | -1.255832 |

In [86]:

```
hier_df.groupby(level='cty', axis=1).describe()
```

Out[86]:

|     |     |       | count | mean      | std      | min       | 25%       | 50%       | 75%       | max       |
|-----|-----|-------|-------|-----------|----------|-----------|-----------|-----------|-----------|-----------|
| cty | cty | tenor |       |           |          |           |           |           |           |           |
| JP  | JP  | 1     | 4.0   | -0.750468 | 0.593964 | -1.479527 | -1.112206 | -0.637598 | -0.275860 | -0.247150 |
|     |     | 3     | 4.0   | -0.097629 | 0.364646 | -0.529011 | -0.332976 | -0.051892 | 0.183455  | 0.242282  |
| US  | US  | 1     | 4.0   | 0.438763  | 0.783844 | -0.574041 | 0.030656  | 0.577680  | 0.985787  | 1.173735  |
|     |     | 3     | 4.0   | -0.282242 | 0.501655 | -1.022996 | -0.375467 | -0.071079 | 0.022146  | 0.036186  |
|     |     | 5     | 4.0   | 0.356805  | 1.172408 | -1.255832 | 0.045183  | 0.567655  | 0.879277  | 1.547744  |



In [87]:

```
hier_df.groupby(level='tenor', axis=1).describe()
```

Out[87]:

|       |     |       | count | mean      | std      | min       | 25%       | 50%       | 75%       | m       |
|-------|-----|-------|-------|-----------|----------|-----------|-----------|-----------|-----------|---------|
| tenor | cty | tenor |       |           |          |           |           |           |           |         |
| 1     | US  | 1     | 4.0   | 0.438763  | 0.783844 | -0.574041 | 0.030656  | 0.577680  | 0.985787  | 1.1737  |
|       | JP  | 1     | 4.0   | -0.750468 | 0.593964 | -1.479527 | -1.112206 | -0.637598 | -0.275860 | -0.2471 |
| 3     | US  | 3     | 4.0   | -0.282242 | 0.501655 | -1.022996 | -0.375467 | -0.071079 | 0.022146  | 0.0361  |
|       | JP  | 3     | 4.0   | -0.097629 | 0.364646 | -0.529011 | -0.332976 | -0.051892 | 0.183455  | 0.2422  |
| 5     | US  | 5     | 4.0   | 0.356805  | 1.172408 | -1.255832 | 0.045183  | 0.567655  | 0.879277  | 1.5477  |



In [ ]:

```
# Data Aggregation

# Aggregations refer to any data transformation that produces scalar values from arrays.
# The preceding examples have used several of them, including mean, count, min, and sum.
# You may wonder what is going on when you invoke mean() on a GroupBy object.
# Many common aggregations, such as those found in Table 10-1, have optimized implementations.
# However, you are not limited to only this set of methods.

# Optimized groupby methods:

#      Function name      Description
#      count              Number of non-NA values in the group
#      sum                 Sum of non-NA values
#      mean                Mean of non-NA values
#      median              Arithmetic median of non-NA values
#      std, var            Unbiased (n - 1 denominator) standard deviation and variance
#      min, max            Minimum and maximum of non-NA values prod Product of non-NA values
#      first, last         First and last non-NA values

# You can use aggregations of your own devising and additionally call any method that
# is also defined on the grouped object. For example, you might recall that quantile
# computes sample quantiles of a Series or a DataFrame's columns.

# While quantile is not explicitly implemented for GroupBy, it is a Series method and
# thus available for use. Internally, GroupBy efficiently slices up the Series,
# calls piece.quantile(0.9) for each piece,
# and then assembles those results together into the result object.
```

In [88]:

```
df
```

Out[88]:

|   | key1 | key2 | data1     | data2     |
|---|------|------|-----------|-----------|
| 0 | a    | one  | -0.699511 | -0.388558 |
| 1 | a    | two  | -0.345863 | 0.535094  |
| 2 | b    | one  | -0.749546 | -1.071714 |
| 3 | b    | two  | -2.196750 | 0.524351  |
| 4 | a    | one  | 0.264189  | 2.621417  |

In [91]:

```
grouped = df.groupby('key1')  
grouped
```

Out[91]:

<pandas.core.groupby.generic.DataFrameGroupBy object at 0x00000039F2156C08>

In [92]:

```
grouped['data1'].quantile(0.9)
```

Out[92]:

```
key1  
a    0.142179  
b   -0.894266  
Name: data1, dtype: float64
```

In [98]:

```
# Quantile and Bucket Analysis

# Pandas has some tools, in particular cut and qcut,
# for slicing data up into buckets with bins of your choosing or by sample quantiles.
# Combining these functions with groupby makes it convenient to perform bucket
# or quantile analysis on a dataset.

# Consider a simple random dataset and an equal-length bucket categorization using cut:

frame = pd.DataFrame({'data1': np.random.randn(1000),
                      'data2': np.random.randn(1000)})
frame
```

Out[98]:

|     | data1     | data2     |
|-----|-----------|-----------|
| 0   | -0.715645 | 0.823377  |
| 1   | 0.715335  | -1.178312 |
| 2   | -1.191247 | 0.205116  |
| 3   | 0.059871  | 1.524632  |
| 4   | 0.473645  | 0.742339  |
| ... | ...       | ...       |
| 995 | 0.003551  | 0.153862  |
| 996 | -0.001380 | 0.987662  |
| 997 | -0.193659 | -0.668325 |
| 998 | -1.016595 | -0.978162 |
| 999 | -2.840384 | -0.171102 |

1000 rows × 2 columns

In [99]:

```
print(frame)
```

```
      data1    data2
0  -0.715645  0.823377
1   0.715335 -1.178312
2  -1.191247  0.205116
3   0.059871  1.524632
4   0.473645  0.742339
..      ...      ...
995  0.003551  0.153862
996 -0.001380  0.987662
997 -0.193659 -0.668325
998 -1.016595 -0.978162
999 -2.840384 -0.171102
```

[1000 rows x 2 columns]

In [100]:

```
frame.unstack()
```

Out[100]:

```
data1  0    -0.715645
       1     0.715335
       2    -1.191247
       3     0.059871
       4     0.473645
```

```
data2  995     0.153862
       996     0.987662
       997    -0.668325
       998    -0.978162
       999    -0.171102
```

Length: 2000, dtype: float64


In [102]:

```
frame.T
```

Out[102]:

|              | 0         | 1         | 2         | 3        | 4        | 5         | 6         | 7         |       |
|--------------|-----------|-----------|-----------|----------|----------|-----------|-----------|-----------|-------|
| <b>data1</b> | -0.715645 | 0.715335  | -1.191247 | 0.059871 | 0.473645 | -0.776308 | -0.048923 | -0.714903 | -2.33 |
| <b>data2</b> | 0.823377  | -1.178312 | 0.205116  | 1.524632 | 0.742339 | 1.582948  | -0.755365 | -0.086841 | 0.49  |

2 rows x 1000 columns



In [104]:

```
quartiles = pd.cut(frame.data1, 4)
quartiles
```

Out[104]:

```
0      (-1.294, 0.253]
1      (0.253, 1.8]
2      (-1.294, 0.253]
3      (-1.294, 0.253]
4      (0.253, 1.8]
...
995    (-1.294, 0.253]
996    (-1.294, 0.253]
997    (-1.294, 0.253]
998    (-1.294, 0.253]
999    (-2.847, -1.294]
Name: data1, Length: 1000, dtype: category
Categories (4, interval[float64]): [(-2.847, -1.294] < (-1.294, 0.253] < (0.253, 1.8] < (1.8, 3.347]]
```

In [106]:

```
quartiles.head(10)
```

Out[106]:

```
0      (-1.294, 0.253]
1      (0.253, 1.8]
2      (-1.294, 0.253]
3      (-1.294, 0.253]
4      (0.253, 1.8]
5      (-1.294, 0.253]
6      (-1.294, 0.253]
7      (-1.294, 0.253]
8      (-2.847, -1.294]
9      (0.253, 1.8]
Name: data1, dtype: category
Categories (4, interval[float64]): [(-2.847, -1.294] < (-1.294, 0.253] < (0.253, 1.8] < (1.8, 3.347]]
```

In [108]:

```
quartiles[:10]
```

Out[108]:

```
0      (-1.294, 0.253]
1      (0.253, 1.8]
2      (-1.294, 0.253]
3      (-1.294, 0.253]
4      (0.253, 1.8]
5      (-1.294, 0.253]
6      (-1.294, 0.253]
7      (-1.294, 0.253]
8      (-2.847, -1.294]
9      (0.253, 1.8]
```

Name: data1, dtype: category

Categories (4, interval[float64]): [(-2.847, -1.294] < (-1.294, 0.253] < (0.253, 1.8] < (1.8, 3.347]]

In [110]:

```
# The Categorical object returned by cut can be passed directly to groupby.
# So we could compute a set of statistics for the data2 column like so:
```

```
grouped = frame.data2.groupby(quartiles)
grouped
```

Out[110]:

```
<pandas.core.groupby.generic.SeriesGroupBy object at 0x00000039F21CCAC8>
```



In [112]:

```
# These were equal-length buckets; to compute equal-size buckets based on sample quantile s, use qcut.
# I'll pass labels=False to just get quantile numbers:

# Return quantile numbers

grouping = pd.qcut(frame.data1, 10, labels=False)
grouping
```

Out[112]:

```
0      2
1      7
2      1
3      5
4      6
..
995    5
996    5
997    4
998    1
999    0
Name: data1, Length: 1000, dtype: int64
```

In [118]:

```
# Filling Missing Values with Group-Specific Values

# When cleaning up missing data, in some cases you will replace data observations using dropna,
# but in others you may want to impute (fill in) the null (NA) values using a fixed value
# or some value derived from the data. fillna is the right tool to use;
# for example, here I fill in NA values with the mean:

s = pd.Series(np.random.randn(6))
s
```

Out[118]:

```
0    -0.562639
1    -1.130607
2    -0.069880
3     1.280894
4     0.224934
5    -0.700833
dtype: float64
```

In [119]:

```
s[::2] = np.nan # filling s with Nan values at intervals of 2  
s
```

Out[119]:

```
0      NaN  
1  -1.130607  
2      NaN  
3   1.280894  
4      NaN  
5  -0.700833  
dtype: float64
```

In [120]:

```
s.fillna(s.mean())
```

Out[120]:

```
0  -0.183515  
1  -1.130607  
2  -0.183515  
3   1.280894  
4  -0.183515  
5  -0.700833  
dtype: float64
```

In [127]:

```
# Making the fill value to vary by group

# Suppose you need the fill value to vary by group. One way to do this is to group the
# data and use apply with a function that calls fillna on each data chunk.

# Here is some sample data on US states divided into eastern and western regions:

states = ['Ohio', 'New York', 'Vermont', 'Florida', 'Oregon', 'Nevada', 'California', 'Idaho']

group_key = ['East'] * 4 + ['West'] * 4

data = pd.Series(np.random.randn(8), index=states)

print(data)
print(group_key)
```

```
Ohio          0.105849
New York     -1.428066
Vermont      -0.453576
Florida     -0.124419
Oregon        0.493002
Nevada        0.361317
California   -1.791414
Idaho         2.516672
dtype: float64
['East', 'East', 'East', 'East', 'West', 'West', 'West', 'West']
```

In [123]:

```
# Let's set some values in the data to be missing:

data[['Vermont', 'Nevada', 'Idaho']] = np.nan
data
```

Out[123]:

```
Ohio          0.304550
New York      0.872126
Vermont         NaN
Florida     -0.528893
Oregon        0.760108
Nevada         NaN
California    0.069828
Idaho         NaN
dtype: float64
```

In [124]:

```
data.groupby(group_key).mean()
```

Out[124]:

```
East    0.215928
West    0.414968
dtype: float64
```

In [129]:

```
# We can fill the NA values using the group means like so:
```

```
fill_mean = lambda g: g.fillna(g.mean())
data.groupby(group_key).apply(fill_mean)
```

Out[129]:

```
Ohio          0.105849
New York      -1.428066
Vermont       -0.453576
Florida       -0.124419
Oregon         0.493002
Nevada         0.361317
California    -1.791414
Idaho          2.516672
dtype: float64
```

In [130]:

```
# Using predefined fill values that vary by group
```

```
# In another case, you might have predefined fill values in your code that vary by group.  
# Since the groups have a name attribute set internally, we can use that:
```

```
fill_values = {'East': 0.5, 'West': -1}  
fill_func = lambda g: g.fillna(fill_values[g.name])  
data.groupby(group_key).apply(fill_func)
```

Out[130]:

```
Ohio          0.105849
New York      -1.428066
Vermont       -0.453576
Florida       -0.124419
Oregon         0.493002
Nevada         0.361317
California    -1.791414
Idaho          2.516672
dtype: float64
```

In [132]:

```
# Random Sampling and Permutation

# Drawing a random sample (with or without replacement) from a large dataset,
# eg: constructing a deck of English-style playing cards

# Suppose you wanted to draw a random sample (with or without replacement) from a large dataset
# for Monte Carlo simulation purposes or some other application.
# There are a number of ways to perform the “draws”; here we use the sample method for Series.

# Hearts, Spades, Clubs, Diamonds
suits = ['H', 'S', 'C', 'D']
card_val = (list(range(1, 11)) + [10] * 3) * 4
base_names = ['A'] + list(range(2, 11)) + ['J', 'K', 'Q']
cards = []
for suit in ['H', 'S', 'C', 'D']:
    cards.extend(str(num) + suit for num in base_names)

deck = pd.Series(card_val, index=cards)
deck
```

Out[132]:

|     |    |
|-----|----|
| AH  | 1  |
| 2H  | 2  |
| 3H  | 3  |
| 4H  | 4  |
| 5H  | 5  |
| 6H  | 6  |
| 7H  | 7  |
| 8H  | 8  |
| 9H  | 9  |
| 10H | 10 |
| JH  | 10 |
| KH  | 10 |
| QH  | 10 |
| AS  | 1  |
| 2S  | 2  |
| 3S  | 3  |
| 4S  | 4  |
| 5S  | 5  |
| 6S  | 6  |
| 7S  | 7  |
| 8S  | 8  |
| 9S  | 9  |
| 10S | 10 |
| JS  | 10 |
| KS  | 10 |
| QS  | 10 |
| AC  | 1  |
| 2C  | 2  |
| 3C  | 3  |
| 4C  | 4  |
| 5C  | 5  |
| 6C  | 6  |
| 7C  | 7  |
| 8C  | 8  |
| 9C  | 9  |
| 10C | 10 |
| JC  | 10 |
| KC  | 10 |
| QC  | 10 |
| AD  | 1  |
| 2D  | 2  |
| 3D  | 3  |
| 4D  | 4  |
| 5D  | 5  |
| 6D  | 6  |
| 7D  | 7  |
| 8D  | 8  |
| 9D  | 9  |
| 10D | 10 |
| JD  | 10 |
| KD  | 10 |
| QD  | 10 |

dtype: int64

In [133]:

```
# So now we have a Series of Length 52 whose index contains card names and values are the ones used  
# in Blackjack and other games (to keep things simple, I just let the ace 'A' be 1):  
deck[:13]
```

Out[133]:

```
AH      1  
2H      2  
3H      3  
4H      4  
5H      5  
6H      6  
7H      7  
8H      8  
9H      9  
10H     10  
JH      10  
KH      10  
QH      10  
dtype: int64
```

In [134]:

```
# Drawing a hand of five cards from the deck could be written as:  
  
def draw(deck, n=5):  
    return deck.sample(n)  
  
draw(deck)
```

Out[134]:

```
3C      3  
JS      10  
JC      10  
JD      10  
8H      8  
dtype: int64
```

In [136]:

```
# Drawing two random cards from each suit

# Suppose you wanted two random cards from each suit.
# Because the suit is the last character of each card name, we can group based on this
and use apply:

get_suit = lambda card: card[-1] # last letter is suit
deck.groupby(get_suit).apply(draw, n=2)
```

Out[136]:

```
C   7C      7
    5C      5
D  10D     10
    9D      9
H   7H      7
    5H      5
S   AS      1
    KS     10
dtype: int64
```

In [137]:

```
# alternatively, we could write:

deck.groupby(get_suit, group_keys=False).apply(draw, n=2)
```

Out[137]:

```
8C      8
7C      7
JD     10
2D      2
JH     10
10H    10
8S      8
6S      6
dtype: int64
```



In [144]:

```
# Group Weighted Average and Correlation

# Under the split-apply-combine paradigm of groupby, operations between columns in a DataF
rame or two Series,
    # such as a group weighted average, are possible.

# As an example, take this dataset containing group keys, values, and some weights:

df = pd.DataFrame({'category': ['a', 'a', 'a', 'a', 'b', 'b', 'b', 'b'],
                    'data': np.random.randn(8),
                    'weights': np.random.rand(8)})

df
```

Out[144]:

|   | category | data      | weights  |
|---|----------|-----------|----------|
| 0 | a        | 0.108480  | 0.278910 |
| 1 | a        | -0.924634 | 0.059491 |
| 2 | a        | -0.639827 | 0.341055 |
| 3 | a        | -0.271886 | 0.732100 |
| 4 | b        | 0.906096  | 0.881823 |
| 5 | b        | -0.313899 | 0.018645 |
| 6 | b        | 1.566793  | 0.745821 |
| 7 | b        | 0.003243  | 0.425661 |

In [145]:

```
# The group weighted average by category would then be:

grouped = df.groupby('category')
get_wavg = lambda g: np.average(g['data'], weights=g['weights'])
grouped.apply(get_wavg)
```

Out[145]:

```
category
a    -0.313141
b     0.947461
dtype: float64
```

In [156]:

```
data
```

Out[156]:

```
Ohio      0.105849
New York  -1.428066
Vermont   -0.453576
Florida   -0.124419
Oregon     0.493002
Nevada     0.361317
California -1.791414
Idaho      2.516672
dtype: float64
```

In [164]:

```
# Cross-Tabulations

# A cross-tabulation (or crosstab for short) is a special case of a pivot table
# that computes group frequencies.

# Here is an example:

data2 = pd.DataFrame({'Sample': ['1', '2', '3', '4', '5', '6', '7', '8', '9', '10'],
                      'Nationality': ['USA', 'Japan', 'USA', 'Japan', 'Japan', 'Japan', 'USA',
                                     'USA', 'Japan', 'USA'],
                      'Handedness': ['Right-handed', 'Left-handed', 'Right-handed', 'Right-ha
                                     nded', 'Left-handed',
                                     'Right-handed', 'Right-handed', 'Left-handed', 'Right-ha
                                     nded', 'Right-handed']})
data2
```

Out[164]:

|   | Sample | Nationality | Handedness   |
|---|--------|-------------|--------------|
| 0 | 1      | USA         | Right-handed |
| 1 | 2      | Japan       | Left-handed  |
| 2 | 3      | USA         | Right-handed |
| 3 | 4      | Japan       | Right-handed |
| 4 | 5      | Japan       | Left-handed  |
| 5 | 6      | Japan       | Right-handed |
| 6 | 7      | USA         | Right-handed |
| 7 | 8      | USA         | Left-handed  |
| 8 | 9      | Japan       | Right-handed |
| 9 | 10     | USA         | Right-handed |

In [165]:

```
# As part of some survey analysis, we might want to summarize this data by nationality and handedness.
```

```
# You could use pivot_table to do this, but the pandas.crosstab function can be more convenient:
```

```
pd.crosstab(data2.Nationality, data2.Handedness, margins=True) #cross_tab
```

Out[165]:

| Handedness  | Left-handed | Right-handed | All |
|-------------|-------------|--------------|-----|
| Nationality |             |              |     |
| Japan       | 2           | 3            | 5   |
| USA         | 1           | 4            | 5   |
| All         | 3           | 7            | 10  |

## The End